

Face Biometrics under Spoofing Attacks: Vulnerabilities, Countermeasures, Open Issues and Research Directions

Abdenour Hadid

Center for Machine Vision Research (CMV), University of Oulu, Finland

Abstract—Among tangible threats and vulnerabilities facing current biometric systems are spoofing attacks. A spoofing attack occurs when a person tries to masquerade as someone else by falsifying data and thereby gaining illegitimate access and advantages. Recently, an increasing attention has been given to this research problem. This can be attested by the growing number of articles and the various competitions that appear in major biometric forums. We have recently participated in a large consortium (TABULARASA) dealing with the vulnerabilities of existing biometric systems to spoofing attacks with the aim of assessing the impact of spoofing attacks, proposing new countermeasures, setting standards/protocols, and recording databases for the analysis of spoofing attacks to a wide range of biometrics including face, voice, gait, fingerprints, retina, iris, vein, electro-physiological signals (EEG and ECG). The goal of this position paper is to share the lessons learned about spoofing and anti-spoofing in face biometrics, and to highlight open issues and future directions.

I. INTRODUCTION

Identity management using biometrics has nowadays become a reality mainly because of the biometric passports (e-passports) and also because of the presence of more and more biometric enabled-applications for personal computers. However, despite the significant progress in the recent decades [11], biometric systems are, unfortunately, vulnerable to attacks. A spoofing attack occurs when a person tries to masquerade as someone else by falsifying data and thereby gaining illegitimate access and advantages. This is currently a major problem for companies willing to market information security solutions based on biometric authentication technologies. For instance, some laptops of Lenovo, Asus and Toshiba come with built-in webcams and embedded biometric systems that authenticate users by scanning their faces. However, in 2009, the Security and Vulnerability Research Team of the University of Hanoi (Vietnam) demonstrated at Black Hat 2009 conference, the world's premier technical security conference, how to easily spoof and bypass these systems (Lenovo's Veriface III, Asus' SmartLogon V1.0.0005, and Toshiba's Face Recognition 2.0.2.32 - each set to its highest security level) using fake facial images of the legitimate user, thus gaining access to the laptops. This vulnerability is now listed in the National Vulnerability Database of the National Institute of Standards and Technology (NIST) in the US. More recently (September 2013), Apple Inc. released its new device, iPhone 5s - with a Touch ID fingerprint sensor for log-in and making users' data more secure. Less than two days later, a German

hacker collective, Chaos Computer Club, claimed and then demonstrated the spoofing of the iPhone 5S with a gummy finger. These two examples, among several others, highlight tangible threats and vulnerabilities in current biometric-based information security. Thus, there is an urgent need for efficient and reliable solutions for detecting and circumventing spoofing attacks. The typical countermeasure to a spoofing attack is liveness detection that aims at detecting some physiological signs of life. It is also assumed that multi-modal systems (e.g. combining face and voice biometric modalities) are in principle more difficult to spoof than uni-modal systems. Thus, gait, face and iris verification could also be performed jointly. However, preliminary investigations indicate that spoofing only one modality can actually be enough to weaken the fusion rule and crack a biometric system protected by multiple modalities.

We have recently participated in a large consortium (TABULARASA EU project, 2010-2014), dealing with the vulnerabilities of existing biometric systems to spoofing attacks with the aim of assessing the impact of spoofing attacks, proposing new countermeasures, setting standards/protocols, and recording databases for the analysis of spoofing attacks to a wide range of biometrics including face, voice, gait, fingerprints, retina, iris, vein, electro-physiological signals (EEG and ECG). The goal of this position paper is to discuss the lessons learned about spoofing and anti-spoofing in face biometrics, and to highlight open issues and future directions.

The rest of this paper is organized as follows. In the next section, Section II, we present an experimental analysis demonstrating the vulnerability of face biometrics to spoofing attacks. Existing databases for studying face anti-spoofing are then described in Section III, whereas Section IV reviews some proposed methods in the literature to cope with face spoofing attacks. Some open issues and future directions are discussed in Section V. Finally, a concluding remarks are drawn in Section VI.

II. VULNERABILITY OF FACE BIOMETRICS TO SPOOFING ATTACKS

To gain insight into the vulnerabilities of face biometric systems when confronted to spoofing attacks, we experimentally analyzed the performance of a baseline system, not trained to handle spoofing attacks, on a challenging 2D face spoofing database known as the REPLAY-ATTACK database [8]. The

chosen baseline face verification system is developed by IDIAP research institute (Switzerland) and uses a part-based face representation and Gaussian mixture models (GMMs) [19]. We briefly describe below the baseline face verification system, the spoofing attack database, the experimental setup and importantly the obtained results which clearly assess the significant vulnerabilities of face biometrics to spoofing attacks.

A. The Baseline Face Biometric System

The face verification system, proposed by McCool and Marcel in [19], is chosen as the baseline system for face authentication. The system combines a part-based face representation and Gaussian mixture models (GMMs). The system divides the face into blocks, and treats each block as a separate observation of the same underlying signal (the face). A feature vector is thus obtained from each block by applying the Discrete Cosine Transform (DCT). The distribution of the feature vectors is then modeled using GMMs.

For feature extraction, the face is normalized, registered and cropped. This cropped and normalized face is divided into blocks (parts) and from each block (part) a feature vector is obtained. Each feature vector is treated as a separate observation of the same underlying signal (in this case the face) and the distribution of the feature vectors is modeled using GMMs. The feature vectors from each block are obtained by applying the DCT [25].

Once the feature vectors are calculated, feature distribution modelling is achieved by performing background model adaptation of GMMs [6], [18]. Background model adaptation first involves the training a world (background) model Ω_{world} from a set of faces and then the derivation of client models Ω_{client}^i for client i by adapting the world model to match the observations of the client. The adaptation is performed using a technique called mean only adaptation [26].

To verify an observation, x , it is scored against both the client (Ω_{client}^i) and world (Ω_{model}) model. The two models, Ω_{client}^i and Ω_{world} , produce a log-likelihood score which is then combined using the log-likelihood ratio (LLR) to produce a single score. This score is used to assign the observation to the world class of faces (not the client) or the client class of faces (it is the client) based on a predefined threshold τ .

B. The Face Spoofing Attack Database

To analyze the performance of the face baseline system under spoofing attacks, we considered the REPLAY-ATTACK face spoofing database [8] which consists of 1300 video clips comprising of real-accesses or photo and video attack attempts to different 50 identities, under different lighting conditions. The data is split into 4 sub-groups comprising enrollment, training, development and test data. Clients that appear in one of the last three data sets do not appear in any other set, while the enrollment set includes all clients.

All videos are generated by either having a (real) client trying to access a laptop through a built-in webcam or by displaying a photo or a video recording of the same client for

at least 9 seconds. In total, 20 attack videos were recorded for each client and 6 videos were captured for real accesses yielding in:

- **Enrollment set:** containing 100 videos (2 per client) for exclusively studying the baseline performance of face recognition systems;
- **Training set:** containing 60 real-accesses and 300 attacks;
- **Development set:** containing 60 real-accesses and 300 attacks;
- **Test set:** containing 80 real-accesses and 400 attacks.

Examples of real accesses and attacks from the REPLAY-ATTACK database are shown in Figure 1. The full description of the database and its associated protocol can be found in [8].



Fig. 1. Examples of real accesses and attacks. In the top row, samples from controlled scenario. In the bottom row, samples from adverse scenario. Columns from left to right show examples of real access, printed photograph, mobile phone and tablet attacks.

C. Experimental Setup

We started by analyzing how the baseline face system behaves on the database's licit protocol setup (i.e. only using real-access attempts). In this mode, the system is tested for how well it can recognize real users authenticating against their templates and how well it can reject real users authenticating against other users' templates (i.e. imposters). The performance is measured objectively by observing the rate of users rejected when authenticating against their own template (False Rejection Rate) and by the rate of users accepted when authenticating against someone else's template (False Acceptance Rate). In this way, we establish the baseline recognition performance of the baseline face recognition for licit access (enrollment attempts and authentication tries).

To determine the baseline face recognition performance, we computed scores exhaustively for all videos from the development and test set real-accesses, probing for every identity in the set against all other models in the same set, **without** intermixing across development and test sets. The scores generated from matched client videos and models within the subset are considered true client accesses and contribute to the licit Acceptance Rate, while all others, imposters, contributing to the licit Rejection Rate. By varying the classification threshold on the test set, we obtain the recognition performance for the system. In this context, the FAR (False Acceptance Rate) is

considered as the rate of impostors that are wrongly classified by the system as true-claimants. The FRR (False Rejection Rate) is the rate of true claimants that the system falsely classified as impostors.

To determine the robustness of the baseline system when exposed to spoofing attacks, we keep the models as trained during the licit protocol performance assessment and try attacks to the models with matching identity. A successful attack is accomplished when the system confuses a spoofing attempt with the corresponding matched user template. In this mode, the FAR corresponds to the rate of attacks that are accepted by the system when spoofed. The FRR corresponds to the rate of real-access attempts that are incorrectly dismissed by the system as attacks.

D. Experimental Results

The results of the experiments are presented in Figure 2 in terms of detection error trade-off (DET) profiles which illustrate the dynamic behavior of a biometric system as the decision threshold is changed, i.e. how the false acceptance rate varies according to the false rejection rate. We also show in Figure 3 the score distributions of true claimants, impostors, and spoofing attacks. By comparing these three distributions, we can observe how spoofed data is closer to information from true claimants than non-spoofed data from an average impostor trying to access the system illegally.




Fig. 2. DET curves for face baseline verification and spoofing attacks. Note that the plot of the baseline verification performance is not visible because the equal error rate is 0 (perfect verification).

From these results, it is possible to notice that the face baseline system achieves perfect performance when not confronted to spoofing attacks, yielding in an easily separable score distribution clusters of the impostors and true claimants (Figure 3) and an equal error rate (EER) of 0% in the DET




Fig. 3. Face Score Distribution

curve (Figure 2). This result indicates that the chosen baseline system is performing very well for this database.

Under spoofing attacks, the scenario is however quite different. The system tends to suffer in a significant way as the score distributions of the true claimants and the attacks overlap as shown in Figure 3. The dashed curve displayed in Figure 2 shows what happens with the FAR as attacks are introduced. Because there are two different sets of false acceptances on the plot, the two FAR curves are overloaded: The solid line shows the FAR for impostors whereas the dashed line shows the FAR for attacks. The way to read such plot is as follows: first one should locate on the horizontal axis an operating point for the FRR and then read both the baseline performance (solid line) and the performance if attacks are introduced (dashed line). For example, if we take an FRR of 0.1%, the baseline system gives a FAR of 0%. If attacks are introduced though, the FAR will increase to above 80%. This means that, for an FRR of 0.1%, the 2D face spoofing database is 80% effective in spoofing the system. The curves in Figure 2 hence indicate that the baseline system is very sensitive to spoofing attacks and this is the case not only at FRR of 0.1% but for any chosen FRR value.

These experimental results show that the face baseline system achieves perfect performance when not confronted to spoof attacks. The performance sharply degrades in the presence of spoofing attacks. These results exemplify the vulnerability of face biometrics against spoofing attacks. These findings are used as motivations for developing countermeasures.

III. EXISTING SPOOFING ATTACK DATABASES

NUAA Photo Imposter Database is among the first public datasets for studying anti-spoofing in face recognition. It was released in 2010, accompanying the work of Tan et al. in [28] in which the authors explored the Lambertian reflectance model to encode the differences between the 2D images of the face presented during an attack and a real (3D) face shown

in real-access attempts. Following the trend of similar past work [16], [4], the authors focused on the binary classification task of face spoofing detection considering pictures of real-accesses and attacks recorded with a conventional webcam. NUA Photo Imposter Database is publicly available and is mainly useful for studying texture-based approaches to spoofing detection.

As shown by Anjos et al. [2], [1], techniques for anti-spoofing can also exploit motion artifacts present in attacks to discriminate spoofing attempts. In [2], the authors made available a public dataset composed of printed photograph attacks and real-accesses, in which the samples available for the training and evaluating spoofing classifiers are videos. The PRINT-ATTACK database can be used to devise anti-spoofing methods based on texture, motion or both [7]. An extension of this database, called the PHOTO-ATTACK database, providing photo attacks using different attack media such as mobile phones and tablets was introduced in [1]. Another extension called REPLAY-ATTACK database, also bringing video attacks using mobile phones and tablets was introduced in [8] and used in the experiments in the previous section.

Zhang et al. have also recorded and released a public dataset for face anti-spoofing containing challenging short video sequence of attacks to 50 different identities using printed photographs and videos displayed through a tablet screen [31]. The photo attacks in this database include warping. The face video attacks in this database can be used for evaluating countermeasures based on motion, texture or both.

Very recently, Erdogmus and Marcel [10] made publicly available a 3D mask database, called 3DMAD, composed of real access and mask attack videos of 17 different subjects recorded by Microsoft Kinect sensor. This database is mainly used for evaluating anti-spoofing measures on 2D face recognition.

IV. FACE ANTI-SPOOFING METHODS

We discuss in this section some existing works in the literature on face anti-spoofing. Short surveys of some schemes against photograph spoofing attacks can be found in [23], [20].

The typical countermeasure to spoofing attacks is liveness detection that aims at detecting physiological signs of life (such as eye blinking, facial expression changes and mouth movements). For instance, Pan et al. [23] proposed an eyeblink-based anti-spoofing method by integrating a structured prediction method whereas Kollreider et al. [12] presented an optical-flow based method to capture the subtle motion of face images. While such countermeasures may work in cases of attacks using photographs, they are generally ineffective when using a video (or simply shaking the photograph before the camera) as a mean of spoofing. Some researchers attempted to counter video spoofing by using structure from motion to calculate the depth information. Again, this may not work in case of spoofing attacks using 3D masks, for instance. Some current face anti-spoofing methods are based on the analysis of the skin properties such as the analysis of skin texture and skin reflectance [23], [20].

In [17], Li et al. described a method for print-attack detection by exploiting differences in the 2D Fourier spectra comparing the hard-copies of client faces and real-accesses. In that work, the authors derive the probability of attack by applying a high-pass filter to the spectra of the sample being analyzed and computing a score which is then classified according to some heuristic. The method works well for down-sampled photos of the attacked identity, but is likely to fail for higher-quality samples. The used dataset is not publicly available.

In [3], the authors proposed a method to detect spoofing attacks using printed photos by analyzing the micro-textures present in the material using a linear SVM classifier to achieve a 2.2% False-Acceptance Rate (FAR) against a 13% False-Rejection Rate (FRR). A major limitation of this method is that the input image needs to be reasonably sharp.

In contrast to the works cited above, the authors in [13], [15] presented a technique to evaluate liveness based on a short sequence of images. The work describes a binary detector that evaluates the trajectories of select parts of the face presented to the input sensor using a simplified optical flow analysis followed by an heuristic classifier. Such a classification scheme achieves an equal-error rate of 0.5% for samples of real-accesses extracted from XM2VTS and attacks produced using hard-copies of those data. The same authors also introduced in [14] a method for fusing scores from different expert systems that observe, concurrently, the 3D face motion scheme introduced on the previous work and liveness properties such as eye-blinks or mouth movements.

The works in [21] and [24] bring a real-time liveness detection specifically against photo-spoofing using (spontaneous) eye-blinks which are supposed to occur once every 2-4 seconds in humans. The system developed uses an undirected conditional random field framework to model the eye-blinking that relaxes the independence assumption of generative modelling and state dependence limitations from hidden Markov modelling. The system is tested on a dataset provided by the authors and was made publicly available. Such a dataset is composed of short video clips of eye-blinks and spoofing attempts using photographs. The attacks are not solely composed of still images but also arbitrary shaking behavior which increases the task difficulty. With this setup, the proposed detector is able to achieve 95.7% true-positive classification against a false alarm of less than 0.1% when considering a simultaneous blink of both eye lids in all test samples. A later work by the same authors [22] augment the number of countermeasures deployed to include a scene context matching that helps preventing video-spoofing in stationary face-recognition systems. To achieve this, the eye-blink detector output scores are fused with the output of a simple local-binary-pattern- χ^2 detector. The scene context detector uses some carefully chosen fiducial points coming from near regions outside the face boundaries that characterize the expected scene context. To test this new setup, the authors constructed a new private dataset with which they obtained an almost perfect scoring - 99.5% true-rejection against 100% true-acceptance.

In [5], Bao et al. proposed a method to detect attacks produced with planar media (such as paper or screens) using motion estimation by optical flow. Movement of planar objects is categorized as translation, rotation, normal or swing and 8 quantities that express the amount of these movements extracted from the analyzed (already) cropped face. The probability of an attack is then computed taking the 8 values and applying them to an *ad-hoc* equation that outputs a single score indicating the probability of a 3D face given the input data. Experiments on a private dataset showed a 6% false-alarm against and 14% false-acceptance in best case.

V. OPEN ISSUES AND FUTURE DIRECTIONS

A. Generalization to Unknown Attacks

Many visual cues for non-intrusive spoofing detection have been already explored and impressive results have been reported on individual databases. However, the varying nature of spoofing attacks and acquisition conditions makes it impossible to predict how single anti-spoofing techniques, e.g. facial texture analysis, can generalize the problem in real-world applications. Moreover, we cannot foresee all possible attack scenarios and cover them in databases because the imagination of the human mind always finds out new tricks to fool existing systems. As one obviously cannot foresee all possible types of fake faces, one-class approach modeling only the genuine facial texture distribution could be a promising direction. This has been successfully applied in voice anti-spoofing [?], for instance.

B. Fusion of Countermeasures

It is reasonable to assume that no single superior technique is able to detect all known, let alone unseen, spoofing attacks. Therefore, the problem of spoofing attacks should be broken down into attack-specific subproblems that are solvable if a proper combination of complementary countermeasures is used. In this manner, a network of attack-specific spoofing detectors could be used to construct a flexible anti-spoofing framework in which new techniques can be easily integrated to patch the existing vulnerabilities in no time when new countermeasures appear. This obviously raises the problem of fusing different spoofing countermeasures which has not been studied much besides the algorithms [27], [29], [30] proposed within the context of the IJCB 2011 competition on counter measures to 2D facial spoofing attacks [7].

C. Biometric System + Countermeasures

A spoofing counter-measure is usually not designated to operate as a stand-alone procedure but in a joint operation with a recognition system. However, most works on anti-spoofing tend to focus only on the spoofing detection part hence omitting to integrate the counter-measure into a recognition system. In practice, integrating the counter-measure will affect the performance of the recognition system. While it will reduce its vulnerability to spoofing attacks, it may also decrease the recognition performance. The open issue is how to combine the spoofing counter-measure and the biometric recognition

so that the combined biometric recognition system is robust to spoofing and does not suffer from reduced recognition accuracy [9].

D. Contextual Information

Face images captured from face spoofs may visually look very similar to the images captured from live faces. Thus, face spoofing detection may be difficult to perform based on only single face image or a relatively short video sequence. Depending on the imaging and fake face quality, it is nearly impossible, even for humans, to tell the difference between a genuine face and a fake one without any scene information or unnatural motion or facial texture patterns. However, we can immediately notice if there is something suspicious in the view, e.g. if someone is holding a video display or a photograph in front of the camera. Therefore, scenic cues can be exploited for determining whether display medium is present in the observed scene.

E. Challenge-Response Approach

Liveness and motion analysis based spoofing detection is rather difficult to perform by observing only spontaneous facial motion during short video sequences. This problem can be simplified by prompting the user to do some specific random action or challenge (such as a smiling and moving the head to the right). The user's response (if any) will provide liveness evidences. This is called challenge-response approach for spoofing detection. The drawback of such an approach is that it requires user cooperation, thus making the authentication process a time-consuming. Another advantage of non-intrusive techniques is that from challenge-response based countermeasures it is rather easy to deduce which liveness cues need to be fooled. For instance, the request for uttering words suggests that analysis of synchronized lip movement and lip reading is utilized, whereas rotating head in a certain direction reveals that the 3D geometry of the head is measured. For non-intrusive approaches, it is usually not known which countermeasures are used, thus the system might be harder to deceive [22].

VI. CONCLUSION

To evaluate the vulnerabilities of face biometric systems when confronted to spoofing attacks, we discussed the performance of a baseline system on a challenging 2D face spoofing database consisting of 1300 video clips of real-accesses and attack attempts to different 50 identities. The 2D face spoofing attack database allows measuring the effectiveness of spoofing attacks or counter-measures to 2D face recognition systems. It is composed of two sets of data: real-accesses and attacks. Real-accesses are used to establish reference performance figures for recognition systems whereas attacks can be used to train spoofing classifiers or measure the impact of spoofing to existing baseline systems. The chosen baseline face verification system uses parts-based Gaussian Mixture Models and provides state-of-the-art performance. The experimental results showed that the face baseline system

achieves perfect performance when not confronted to spoof attacks. The performance sharply degrades in the presence of spoofing attacks. For instance, for a False Rejection Rate (FRR) of 0.1%, the FAR on real impostors goes from 0% to more than 80% when attacks are introduced. These results exemplify the vulnerability of face biometric systems against spoofing attacks.

Without spoofing counter-measures, most of the state-of-the-art facial biometric systems are indeed vulnerable to attacks, since they try to maximize the discriminability between identities without regards to whether the presented trait originates from a living legitimate client or not. The proposed anti-spoofing methods in the literature have shown very encouraging results on individual databases but may lack generalization to varying nature of spoofing attacks that can be encountered in real-world applications. This suggests that a network of attack-specific spoofing detectors maybe needed to tackle different spoofing attacks. The existing databases for spoofing and anti-spoofing analysis have been and are still useful for studying the spoofing problems but one cannot foresee all possible attack scenarios and cover them in databases. As the field evolves, new and more challenging databases can be expected. The imagination of the human mind always finds out new tricks to fool existing biometric systems. As one obviously cannot foresee all possible types of fake faces, one-class approach modeling only the genuine facial texture distribution could be a promising direction.

The open issues and the research directions that have been discussed in this paper are not specific to face biometrics but also hold for other biometric modalities. Due to lack of space, we focused this paper on face biometrics and did not report the results on voice, gait, fingerprints, retina, iris, vein, electro-physiological signals (EEG and ECG). In summary, the investigations in the TABULA RASA EU project showed that most of the biometrics modalities, including the multi-modal combinations, are vulnerable to spoofing attacks with different degrees.

Acknowledgments: The financial support of the Academy of Finland and NOKIA Foundation is acknowledged. This work is done within the FP7 EU European Project TABULARASA.

REFERENCES

- [1] A. Anjos, M. M. Chakka, and S. Marcel. Motion-based counter-measures to photo attacks in face recognition. *IET Biometrics*, 2013.
- [2] A. Anjos and S. Marcel. Counter-measures to photo attacks in face recognition: a public database and a baseline. In *Proceedings of IAPR IEEE International Joint Conference on Biometrics (IJCB)*, Washington DC, USA, 2011.
- [3] J. Bai, T. Ng, X. Gao, and Y. Shi. Is physics-based liveness detection truly possible with a single image? In *Circuits and Systems ISCAS Proceedings of 2010 IEEE International Symposium on*, page 34253428. IEEE, 2010.
- [4] J. Bai, T.-T. Ng, X. Gao, and Y.-Q. Shi. Is physics-based liveness detection truly possible with a single image? In *IEEE International Symposium on Circuits and Systems (ISCAS)*, pages 3425–3428, 2010.
- [5] W. Bao, H. Li, N. Li, and W. Jiang. A liveness detection method for face recognition based on optical flow field. In *2009 International Conference on Image Analysis and Signal Processing*, pages 233–236. IEEE, 2009.
- [6] F. Cardinaux, C. Sanderson, and S. Marcel. Comparison of mlp and gmm classifiers for face verification on XM2VTS. In *Proc. International Conference on Audio- and Video-based Biometric Person Authentication*, pages 911–920, 2003.
- [7] M. M. Chakka, A. Anjos, and S. Marcel. Competition on counter measures to 2-d facial spoofing attacks. In *Proceedings of IAPR IEEE International Joint Conference on Biometrics (IJCB)*, Washington DC, USA, 2011.
- [8] I. Chingovska, A. Anjos, and S. Marcel. On the effectiveness of local binary patterns in face anti-spoofing. In *IEEE BIOSIG 2012*, Sept. 2012.
- [9] I. Chingovska, A. Anjos, and S. Marcel. Anti-spoofing in action: joint operation with a verification system. In *Proceedings of IEEE Conference on Computer Vision and Pattern Recognition, Workshop on Biometrics*, June 2013.
- [10] N. Erdogmus and S. Marcel. Spoofing in 2d face recognition with 3d masks and anti-spoofing with kinect. In *Biometrics: Theory, Applications and Systems Conference (BTAS'13)*, Sept. 2013.
- [11] A. K. Jain, P. Flynn, and A. A. Ross. *Handbook of Biometrics*. Springer, 2008.
- [12] K. Kollreider, H. Fronthaler, and J. Bigun. Non-intrusive liveness detection by face images. *Image and Vision Computing*, 27:233–244, 2009.
- [13] K. Kollreider, H. Fronthaler, and J. Bigun. Evaluating liveness by face images and the structure tensor. In *Fourth IEEE Workshop on Automatic Identification Advanced Technologies AutoID05*, pages 1–6. IEEE, 2005.
- [14] K. Kollreider, H. Fronthaler, and J. Bigun. Verifying liveness by multiple experts in face biometrics. In *2008 IEEE Computer Society Conference on Computer Vision and Pattern Recognition Workshops*, pages 1–6. IEEE, 2008.
- [15] K. Kollreider, H. Fronthaler, and J. Bigun. Non-intrusive liveness detection by face images. *Image and Vision Computing*, 27(3):233–244, 2009.
- [16] J. Li, Y. Wang, T. Tan, and A. K. Jain. Live face detection based on the analysis of fourier spectra. In *Biometric Technology for Human Identification*, pages 296–303, 2004.
- [17] J. Li, Y. Wang, T. Tan, and A. K. Jain. Live face detection based on the analysis of fourier spectra. In *Biometric Technology for Human Identification*, pages 296–303, 2004.
- [18] S. Lucey and T. Chen. A gmm parts based face representation for improved verification through relevance adaptation. In *Computer Vision and Pattern Recognition*, pages 855–861, 2004.
- [19] C. McCool and S. Marcel. Parts-based face verification using local frequency bands. In *Proceedings of IEEE/IAPR International Conference on Biometrics*, 0 2009.
- [20] K. Nixon, V. Aimale, and R. Rowe. Spoof detection schemes. In X, editor, *Handbook of Biometrics*, pages 403–4239. X, 2008.
- [21] G. Pan, L. Sun, Z. Wu, and S. Lao. Eyeblick-based anti-spoofing in face recognition from a generic webcam. *IEEE 11th International Conference on Computer Vision (2007)*, pages 1–8, 2007.
- [22] G. Pan, L. Sun, Z. Wu, and Y. Wang. Monocular camera-based face liveness detection by combining eyeblink and scene context. *Telecommunication Systems*, 47(3-4):215–225, 2011.
- [23] G. Pan, Z. Wu, and L. Sun. Liveness detection for face recognition. In K. Delac, M. Grgic, and M. S. Bartlett, editors, *Recent Advances in Face Recognition*, page Chapter 9. IN-TECH, 2008.
- [24] G. Pan, Z. Wu, and L. Sun. Liveness detection for face recognition. *Recent Advances in Face Recognition*, (December):109–124, 2008.
- [25] K. R. Rao and P. Yip. *Discrete Cosine Transform: Algorithms, Advantages, Applications*. Academic Press, Boston, 1990.
- [26] D. A. Reynolds. Comparison of background normalization methods for text-independent speaker verification. In *International Conference on Acoustics, Speech, and Signal Processing*, pages 963–966, 1997.
- [27] W. R. Schwartz, A. Rocha, and H. Pedrini. Face Spoofing Detection through Partial Least Squares and Low-Level Descriptors. In *Proceedings of IAPR IEEE International Joint Conference on Biometrics (IJCB)*, Washington DC, USA, 2011.
- [28] X. Tan, Y. Li, J. Liu, and L. Jiang. Face liveness detection from a single image with sparse low rank bilinear discriminative model. In *Proceedings of the 11th European conference on Computer vision: Part VI, ECCV'10*, pages 504–517, 2010.
- [29] R. Tronci, D. Muntoni, G. Fadda, M. Pili, N. Sirena, G. Murgia, M. Ristori, and F. Roli. Fusion of multiple clues for photo-attack detection in face recognition systems. In *Proceedings of IAPR IEEE International Joint Conference on Biometrics (IJCB)*, Washington DC, USA, 2011.
- [30] J. Yan, Z. Zhang, Z. Lei, D. Yi, and S. Z. Li. Face liveness detection by exploring multiple scenic clues. In *12th International Conference on Control, Automation, Robotics and Vision, (ICARCV2012)*, 2012.
- [31] Z. Zhang, J. Yan, S. Liu, Z. Lei, D. Yi, and S. Z. Li. A face antispoofing database with diverse attacks. In *5th IAPR International Conference on Biometrics (ICB'12)*, 2012.